

EECS 391 Introduction to Artificial Intelligence

Fall 2018, Programming Assignment 2 ("P2")

Due: Fri Dec 7 before midnight. **Total Points:** 100 + 20 extra credit

In this assignment, you will train a linear classifier on the iris dataset. Like HW5, this will be convenient to do in a language with a good plotting library such as Matlab, Mathematica, or python using matplotlib. Submit both your code and your write-up (which should be a pdf file) via Canvas.

Write-up: Your write-up should answer the questions like in the written assignments. When relevant, include the source code in the write-up if it is sufficiently short (say half a page or less), otherwise refer to the file. For example, Exercise 2a asks you to write code to compute the mean-squared error, so you would just need to show the lines for that function. For reference, see §18.6.3-4 of the textbook.

Exercise 1. Linear decision boundaries

- a) Inspect `irisdata.csv` which is provided with the assignment file on Canvas. Write a program that loads the iris data set and plots the 2nd and 3rd iris classes, similar to the plots shown in lecture on neural networks (i.e. using the petal width and length dimensions). 10 P.
- b) Define a function that computes the output of simple one-layer neural network using a logistic non-linearity (linear classification with logistic regression). Use an output of 0 to indicate the 2nd iris class, and 1 to indicate the 3rd. 5 P.
- c) Write a function that plots the decision boundary for the non-linearity above overlaid on the iris data. Choose a boundary (by setting weight parameters by hand) that roughly separates the two classes. 5 P.
- d) Use a surface plot from a 3D plotting library (e.g. in matlab or using matplotlib in python) to plot the output of your neural network over the input space. This should be similar to fig 18.7c in the textbook. 5 P.
- e) Show the output of your simple classifier using examples from the 2nd and 3rd Iris classes. Choose examples that are unambiguous as well as those that are near the decision boundary. 5 P.

Exercise 2. Neural networks (see §18.7 of the textbook)

- a) Write a program that calculates the mean-squared error of the iris data for the neural network defined above. The function should take three arguments: the data vectors, the parameters defining the neural network, and the pattern classes. 10 P.
- b) Compute the mean squared error for two different settings of the weights (i.e. two different decision boundaries). Like above, select these by hand and choose settings that give large and small errors respectively. Plot both boundaries on the dataset as above. 5 P.
- c) Give a mathematical derivation the gradient of the objective function above with respect to the neural network output. You will have to use the derivative of the logistic function as discussed in class. Use w_0 to represent the bias term. You should show and explain each step. 10 P.
- d) Show how the gradient can be written in both scalar and vector form. 5 P.
- e) Write code that computes the summed gradient for an ensemble of patterns. Illustrate the gradient by showing (i.e. plotting) how the decision boundary changes for a small step. 10 P.

Exercise 3. Learning a decision boundary through optimization

- a) Using your code above, write a program that implements gradient descent to optimize the decision boundary for the iris dataset. 5 P.
- b) In your program, include code that shows the progress in two plots: the first should show the current decision boundary location overlaid on the data; the second should show the *learning curve*, i.e. a plot of the objective function as a function of the iteration (see figure 18.18 of the textbook). 5 P.
- c) Run your code on the iris data set starting from a random setting of the weights. Note: you might need to restrict the degree of randomness so that the initial decision boundary is visible somewhere in the 10 P.

plot. In your writeup, show the two output plots at the initial, middle, and final locations of the decision boundary.

d) Explain how you chose the gradient step size. 5 P.

e) Explain how you chose a stopping criterion. 5 P.

Exercise 4. Extra credit: Using a deep learning toolbox

a) Follow the tutorial provided with the assignment to define the neural network above using one of the toolboxes discussed in the tutorials. +5 P.

b) Generalize the network to perform classification on the full iris dataset, i.e. your network should classify all three iris class and use all four data dimensions. Feel free to generalize the network and report your results. +15 P.